Human-in-the-Loop is Not Enough

The Persistent and Evolving Limits of Human Oversight in High-Stakes AI

About me

Ankur Vatsa

Experience, Expertise and Education:

- Built SaaS-based regulatory reporting systems with ML-powered validation
- Led AI/ML adoption across 180+ applications at global investment banks
- Specialist in AI-powered regulatory reporting, risk analytics, and compliance automation
- Expert in Human-AI collaboration patterns in high-stakes financial environments
- AWS Solutions Architect – Professional, with 17+ years in global investment banking technology
- M.Tech. in Artificial Intelligence and Machine Learning – BITS Pilani

Another look at

- **HITL** "Human in the loop" as a remedy we commonly bank upon
- **Diagnose** HITL failure modes in AI systems
- **Evaluate** the insufficiencies of HITL and safeguards to **Consider**
- **Implement** discriminator agents, monitoring, and governance controls
- **Prioritize** remediation steps using a risk-based roadmap

Human-in-the-Loop (HITL)

- **Common remedy** where intelligent systems hand over control to humans in several key scenarios:

Key scenario	Description
Errors and anomalies observed	AI detects anomalies or unexpected outputs
	System confidence drops below threshold
	Error conditions trigger human intervention
Low risk appetite	High-stakes decisions require human approval
	Regulatory compliance mandates human oversight
	Business-critical processes need human validation

- not just "humans reviewing AI outputs" — it's a structured handover mechanism

HITL Triggers (Continued)

Key scenario	Description
Need for domain expertise	Complex edge cases beyond AI training scope
	Nuanced judgment calls requiring contextual understanding
	Ethical considerations and value-based decisions
Legal/regulatory mandates	Healthcare diagnostics requiring physician approval
	Financial lending decisions with fair lending requirements
	Criminal justice applications with due process rights

- HITL offers defence to AI limitations; but faces issues — *speed, scale, bias, skillsgap*

The Case for Enhanced AI Oversight

- **$460M in 45 minutes**: Knight Capital flash crash (2012)
- **10,000 families wrongfully penalized**: Dutch childcare scandal
- **30% error miss rate**: Hospital audits of AI diagnostic reviews
- **95% false positive rate**: AML alert systems overwhelming investigators

HITL alone is insufficient for high-stakes AI deployment.

Five Failure Modes of HITL

1. **Speed Mismatch**: Microsecond AI vs. second-scale humans *(Knight Capital: 45 minutes, $460M)*
2. **Scale Mismatch**: 10,000+ daily alerts vs. limited reviewers *(AML: 95% false positives)*
3. **Overtrust**: Automation bias hides critical errors *(Hospital study: 30% miss rate)*
4. **Skill Erosion**: "Out-of-the-loop" degraded expertise *(Tesla Autopilot incidents)*
5. **Coordination Gaps**: Unclear handovers and protocols *(Dutch childcare: role confusion)*

Financial Services: When Structure Deceives

The Problem:

- AI extracts regulatory requirements from EMIR, MiFID, CFTC rules
- **Hallucination risk**: LLMs invent non-existent field requirements
- **False confidence**: Well-formatted JSON output appears trustworthy
- **Scale challenge**: Hundreds of fields vs. limited compliance staff

Real Impact:

- Potential fines: >$50M for insufficient swap reporting
- Deutsche Bank: $150M AML penalty for ongoing failures

Healthcare: Life-or-Death Automation Bias

The Evidence:

- Hospital audits: Clinicians miss **30% of AI diagnostic errors**
- Emergency departments: Time pressure prevents thorough AI output validation
- Training gaps: Physicians lack awareness of AI system limitations

Why HITL Fails:

- **Automation bias**: Trusting confident AI assessments of "normal" results
- **Mode confusion**: Unclear guidance on when to override AI recommendations
- **Workload pressure**: No time for careful review during patient surges

Transportation: Split-Second Decisions

NHTSA Findings (2023):

- Multiple fatal Tesla Autopilot incidents
- Drivers had **seconds** to react, weren't ready to take control
- "Out-of-the-loop" problem: Skill erosion from over-reliance

The Core Issue:

- AI operates in **milliseconds**
- Human reaction time: **1-3 seconds**
- Handover complexity: Mode awareness, situational context, skill maintenance

Discriminator Agents: Automated Validation

What they are:

Specialized AI systems that detect errors, hallucinations, or inconsistencies in other AI outputs

Architecture Pattern:

```
Source → Primary AI → Output → Discriminator AI → Flagged Items → Human Review
```

Types:

- **Binary classifiers**: Error/No-error detection
- **Consistency checkers**: Cross-reference with source material
- **Adversarial critics**: Challenge model outputs
- **Provenance verifiers**: Trace data lineage

Concrete Monitoring Metrics

Performance Thresholds:

- **Model drift**: Alert if accuracy drops >5% from baseline
- **Human interception**: Warn if <80% error catch rate
- **Confidence escalation**: Review if AI confidence <70% on critical outputs
- **Anomaly rate**: Flag if >10% of outputs marked unusual in 24h

Real Examples:

- **GLUE/SuperGLUE**: NLP benchmarks, target >90% accuracy
- **ImageNet**: Computer vision, target >95% top-5 accuracy
- **Financial datasets**: <1% hallucination rate in regulatory extraction

12-Month Implementation Roadmap

Phase 1: Foundation (0-3 months)

✓ Establish baselines with benchmark datasets

✓ Implement structured logging and audit trails

✓ Define concrete alerting thresholds

Phase 2: Detection (3-6 months)

✓ Deploy continuous drift monitoring

✓ Add discriminator agents for high-risk outputs

✓ Start measuring human interception rates

Phase 3: Prevention (6-12 months)

✓ Implement automated circuit breakers

✓ Add ensemble validation for critical decisions

✓ Establish incident response playbooks

Audit Trail Schema

```
{
  "event_id": "uuid",
  "timestamp": "2025-09-02T10:00:00Z",
  "input": {"source": "doc123"},
  "ai_output": {"model": "v1.2", "prediction": "...", "confidence": 0.85},
  "discriminator": {"model": "discA", "flag": true, "reason": "inconsistency"},
  "human_review": {"reviewer": "user456", "action": "approve", "notes": "..."},
  "outcome": {"final": "approved", "error_detected": false}
}
```

Key Requirements:

- Immutable storage for regulatory compliance
- Accessible to auditors with proper retention policies
- Captures full decision lineage for post-incident analysis

Key Implementation Patterns

Automated Circuit Breakers:

- Financial: Trading halts when anomalies detected (<1 second response)
- Healthcare: Confidence-based escalation to senior clinicians
- Regulatory: Stop processing when hallucination rate exceeds threshold

UI Design for Trust Calibration:

- Show confidence scores and uncertainty ranges
- Highlight source material for AI conclusions
- Force deliberation on high-risk overrides
- Provide "second opinion" views from ensemble models

Summary and Conclusions

HITL is necessary but not sufficient

The evidence:

- Knight Capital: $460M in 45 minutes (speed mismatch)
- Dutch childcare: 10,000 families wrongfully penalized (automation bias)
- Hospital study: 30% diagnostic error miss rate (overtrust)

The solution:
Layered defenses with discriminator agents, automated safeguards, and measurable monitoring

Next steps:
Start with benchmarks and baseline measurements—systems cannot be improved without measurement